

DriverJinn — Downstream Analysis & Poster Plots

Setup

Training Curves

```
KEY = "val_ndcg@50"
fig, ax = plt.subplots(figsize=(8, 6))
for fold, history in sorted(histories.items()):
    if KEY not in history or len(history[KEY]) == 0:
        continue
    y = np.array(history[KEY])
    ax.plot(np.arange(1, len(y) + 1), y,
            label=f"Fold {fold}", color=FOLD_COLORS[fold - 1], alpha=0.85, linewidth=1.5)
ax.set_xlabel("Epoch", fontsize=12)
ax.set_ylabel("NDCG@50", fontsize=12)
ax.set_title("Validation NDCG@50 Across Folds", fontsize=14, fontweight="bold")
ax.legend(loc="best", fontsize=10)
ax.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()
```




Figure 1: Validation NDCG@50 across folds.

```
KEY = "val_auroc"
fig, ax = plt.subplots(figsize=(8, 6))
for fold, history in sorted(histories.items()):
    if KEY not in history or len(history[KEY]) == 0:
        continue
    y = np.array(history[KEY])
    ax.plot(np.arange(1, len(y) + 1), y,
            label=f"Fold {fold}", color=FOLD_COLORS[fold - 1], alpha=0.85, linewidth=1.5)
ax.set_xlabel("Epoch", fontsize=12)
ax.set_ylabel("AUROC", fontsize=12)
ax.set_title("Validation AUROC Across Folds", fontsize=14, fontweight="bold")
ax.legend(loc="lower right", fontsize=10)
ax.yaxis.set_minor_locator(ticker.AutoMinorLocator())
ax.grid(True, alpha=0.3)
plt.tight_layout()
```

```
plt.show()
```




Figure 2: Validation AUROC across folds.

```
rows = []
for fold, history in sorted(histories.items()):
    ndcg = history.get("val_ndcg@50", [])
    auroc = history.get("val_auroc", [])
    rows.append({
        "Fold": fold,
        "Peak NDCG@50": round(max(ndcg), 4) if ndcg else None,
        "Peak AUROC": round(max(auroc), 4) if auroc else None,
        "Epochs": len(ndcg),
    })
df_summary = pd.DataFrame(rows)
mean_row = {"Fold": "Mean",
```

```

        "Peak NDCG@50": round(df_summary["Peak NDCG@50"].mean(), 4),
        "Peak AUROC": round(df_summary["Peak AUROC"].mean(), 4),
        "Epochs": "-"
    }
df_summary = pd.concat([df_summary, pd.DataFrame([mean_row])], ignore_index=True)
df_summary.style.set_caption("Per-fold peak validation metrics")

```

Table 1: Per-fold peak validation metrics

	Fold	Peak NDCG@50	Peak AUROC	Epochs
0	1	0.399200	0.753400	1080
1	2	0.502200	0.751900	680
2	3	0.426200	0.754200	780
3	4	0.524400	0.814400	1200
4	5	0.537500	0.766500	700
5	Mean	0.477900	0.768100	—

Enrichment Curve (Cumulative CGC Precision)

How quickly DriverJinn recovers known cancer driver genes as you move down the ranked list.

```

fig, ax = plt.subplots(figsize=(7, 5))
ks = np.arange(1, len(genes) + 1)
cumulative_t1 = np.cumsum(genes["cgc_tier"] == "T1") / ks
cumulative_any = np.cumsum(genes["cgc_tier"] != "non-CGC") / ks
baseline_t1 = (genes["cgc_tier"] == "T1").sum() / len(genes)
baseline_any = (genes["cgc_tier"] != "non-CGC").sum() / len(genes)


ax.axhline(baseline_t1, color=COLORS["T1"], linestyle="--", linewidth=1.2,
            alpha=0.5, label=f"T1 random baseline ({baseline_t1:.3f})")
ax.axhline(baseline_any, color=COLORS["T2"], linestyle="--", linewidth=1.2,
            alpha=0.5, label=f"Any-CGC random baseline ({baseline_any:.3f})")
ax.plot(ks, cumulative_t1, color=COLORS["T1"], linewidth=2.2, label="CGC Tier 1 precision")
ax.plot(ks, cumulative_any, color=COLORS["T2"], linewidth=2.2, label="Any CGC precision")
for k in [10, 50, 100, 200]:
    if k <= len(genes):
        ax.axvline(k, color="#CCCCCC", linewidth=0.8, linestyle=":")
ax.set_xlim(1, 500)

```

```

ax.set_xlabel("Rank cutoff (K)", fontsize=12)
ax.set_ylabel("Cumulative Precision@K", fontsize=12)
ax.set_title("DriverJinn - Enrichment Curve (GGNet, =0.2)", fontsize=13, fontweight="bold")
ax.legend(fontsize=10)
sns.despine()
plt.tight_layout()
plt.savefig("poster_enrichment_curve.pdf", dpi=300, bbox_inches="tight")
plt.show()

```



Score Distribution by CGC Tier

Separation of DriverJinn scores across CGC Tier 1, Tier 2, and non-CGC genes.

```

fig, ax = plt.subplots(figsize=(6, 5))
order = ["T1", "T2", "non-CGC"]
palette = {k: COLORS[k] for k in order}
sns.violinplot(data=genes, x="cgc_tier", y="mean_score",
               order=order, palette=palette, inner="box", linewidth=1.2, ax=ax)
for tier, color in [("T1", COLORS["T1"]), ("T2", COLORS["T2"])] :
    sub = genes[genes["cgc_tier"] == tier]
    ax.scatter([order.index(tier)] * len(sub), sub["mean_score"],
               color="white", edgecolors=color, s=18, zorder=5, linewidth=0.8)
for i, tier in enumerate(order):
    med = genes.loc[genes["cgc_tier"] == tier, "mean_score"].median()
    ax.text(i, med + 0.005, f"{med:.3f}", ha="center", va="bottom",
           fontsize=9, fontweight="bold", color="black")
ax.set_xlabel("Gene category", fontsize=12)
ax.set_ylabel("DriverJinn mean score", fontsize=12)
ax.set_title("Score Distribution by CGC Status\n(GGNet, =0.2)", fontsize=13, fontweight="bold")
ax.set_xticklabels(["CGC Tier 1", "CGC Tier 2", "Non-CGC"])
sns.despine()
plt.tight_layout()
plt.savefig("poster_score_violin.pdf", dpi=300, bbox_inches="tight")
plt.show()

```


C:\Users\siddh\AppData\Local\Temp\ipykernel_43672\372424057.py:4: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign

```

sns.violinplot(data=genes, x="cgc_tier", y="mean_score",
C:\Users\siddh\AppData\Local\Temp\ipykernel_43672\372424057.py:17: UserWarning: set_ticklabel
ax.set_xticklabels(["CGC Tier 1", "CGC Tier 2", "Non-CGC"])

```



Precision@K Across Folds

Per-fold and mean Precision at multiple rank cutoffs, with error bars.


```
ks = [10, 20, 50, 100, 200]
rows = []
for fold, df in fold_dfs.items():
    df = df.sort_values("rank").reset_index(drop=True)
    df["cgc_t1"] = df["gene_name"].isin(cgc_t1)
    df["cgc_any"] = df["gene_name"].isin(cgc_t1 | cgc_t2)
    for k in ks:
        top_k = df.head(k)
        rows.append({"fold": fold, "K": k,
```

```

        "Precision@K (T1)":      top_k["cgc_t1"].mean(),
        "Precision@K (Any CGC)": top_k["cgc_any"].mean()})
prec_df = pd.DataFrame(rows)
summary = (prec_df.groupby("K")["Precision@K (T1)", "Precision@K (Any CGC)"]
           .agg(["mean", "std"]).reset_index())
summary.columns = ["K", "t1_mean", "t1_std", "any_mean", "any_std"]

fig, ax = plt.subplots(figsize=(7, 5))
width, xs = 0.35, np.arange(len(ks))
ax.bar(xs - width/2, summary["t1_mean"], width, yerr=summary["t1_std"],
       color=COLORS["T1"], alpha=0.85, capsize=4, label="CGC Tier 1",
       error_kw={"linewidth": 1.2})
ax.bar(xs + width/2, summary["any_mean"], width, yerr=summary["any_std"],
       color=COLORS["T2"], alpha=0.85, capsize=4, label="Any CGC",
       error_kw={"linewidth": 1.2})
baseline_t1_v = len(cgc_t1 & set(genes["gene_name"])) / len(genes)
baseline_any_v = len((cgc_t1 | cgc_t2) & set(genes["gene_name"])) / len(genes)
ax.axhline(baseline_t1_v, color=COLORS["T1"], linestyle="--", linewidth=1, alpha=0.5)
ax.axhline(baseline_any_v, color=COLORS["T2"], linestyle="--", linewidth=1, alpha=0.5)
ax.set_xticks(xs)
ax.set_xticklabels([f"@{k}" for k in ks])
ax.set_xlabel("Rank cutoff K", fontsize=12)
ax.set_ylabel("Precision@K (mean ± std, 5 folds)", fontsize=12)
ax.set_title("DriverJinn - Precision@K Across Folds\n(GGNet, =0.2)", fontsize=13, fontweight="bold")
ax.legend(fontsize=10)
sns.despine()
plt.tight_layout()
plt.savefig("poster_precision_at_k.pdf", dpi=300, bbox_inches="tight")
plt.show()

```

Rank Stability Heatmap

For the consensus top-30 genes, how consistent is their rank across all 5 folds?

```
TOP_N      = 30
top_genes = genes.head(TOP_N)["gene_name"].tolist()
rank_matrix = pd.DataFrame(index=top_genes)
for fold, df in fold_dfs.items():
    df = df.sort_values("rank")
    name_to_rank = dict(zip(df["gene_name"], df["rank"]))
    rank_matrix[f"Fold {fold}"] = [name_to_rank.get(g, np.nan) for g in top_genes]

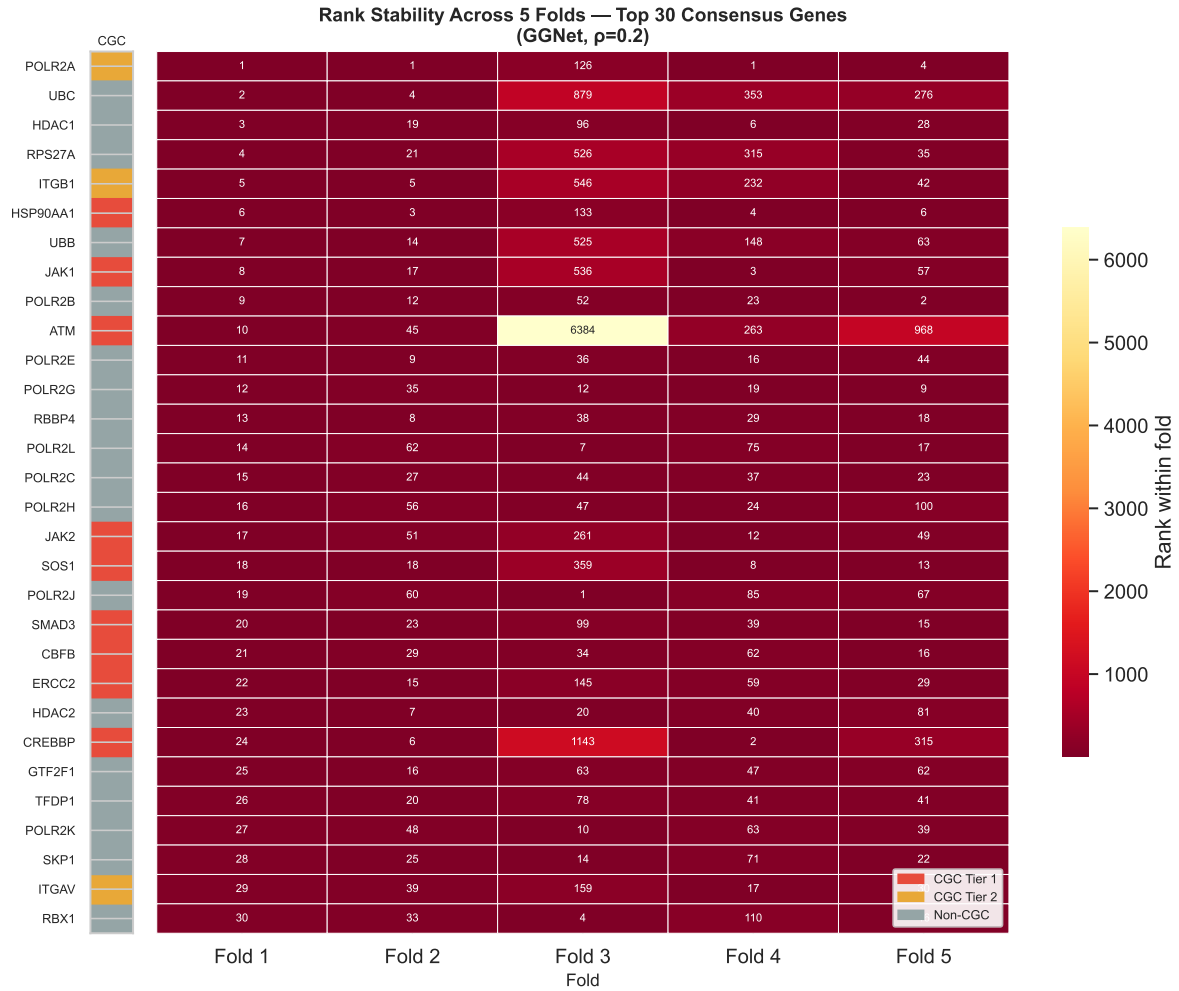
cgc_array = np.array([[
    0 if g in cgc_t1 else 1 if g in cgc_t2 else 2
    for g in top_genes
    ]])
```

```

])).T

fig, axes = plt.subplots(1, 2, figsize=(11, 9),
                        gridspec_kw={"width_ratios": [0.04, 1]})
axes[0].imshow(cgc_array, aspect="auto",
               cmap=plt.cm.colors.ListedColormap([COLORS["T1"], COLORS["T2"], COLORS["non-CGC"]]))
axes[0].set_xticks([])
axes[0].set_yticks(range(TOP_N))
axes[0].set_yticklabels(top_genes, fontsize=8)
axes[0].set_title("CGC", fontsize=8, pad=4)
sns.heatmap(rank_matrix, ax=axes[1], cmap="YlOrRd_r", annot=True, fmt=".0f",
            annot_kws={"size": 7}, linewidths=0.3,
            cbar_kws={"label": "Rank within fold", "shrink": 0.6})
axes[1].set_yticks([])
axes[1].set_xlabel("Fold", fontsize=11)
axes[1].set_title("Rank Stability Across 5 Folds - Top 30 Consensus Genes\n(GGNet, =0.2)",
                  fontsize=12, fontweight="bold")
axes[1].legend(handles=[
    mpatches.Patch(color=COLORS["T1"], label="CGC Tier 1"),
    mpatches.Patch(color=COLORS["T2"], label="CGC Tier 2"),
    mpatches.Patch(color=COLORS["non-CGC"], label="Non-CGC"),
], loc="lower right", fontsize=8, framealpha=0.9)
plt.tight_layout()
plt.savefig("poster_rank_stability.pdf", dpi=300, bbox_inches="tight")
plt.show()

```



Per-Fold Metrics Summary

Overview of $NDCG@50$, $AUROC$, $AUPRC$, $Precision@50$ across all 5 folds.


```
metric_cols = ["NDCG@50", "AUROC", "AUPRC", "Precision@50"]
fold_data   = metrics_folds[["Fold"] + metric_cols].set_index("Fold")
mean_row    = metrics[metrics["Fold"] == "Mean"][metric_cols].values[0]
std_row     = metrics[metrics["Fold"] == "Std"][metric_cols].values[0]

fig, ax = plt.subplots(figsize=(8, 5))
```

```

fold_palette = sns.color_palette("husl", 5)
xs = np.arange(len(metric_cols))
width = 0.13
for i, (fold, row) in enumerate(fold_data.iterrows()):
    ax.bar(xs + (i - 2) * width, row[metric_cols],
           width, alpha=0.65, color=fold_palette[i], label=f"Fold {fold}")
ax.errorbar(xs, mean_row, yerr=std_row, fmt="D", color="black",
            markersize=6, linewidth=1.8, capsize=5, zorder=10, label="Mean ± std")
ax.set_xticks(xs)
ax.set_xticklabels(metric_cols, fontsize=11)
ax.set_ylabel("Score", fontsize=12)
ax.set_ylim(0, 1)
ax.set_title("DriverJinn - Per-Fold Validation Metrics\n(GGNet, =0.2)",
             fontsize=13, fontweight="bold")
ax.legend(fontsize=9, ncol=3, loc="upper right")
sns.despine()
plt.tight_layout()
plt.savefig("poster_metrics_summary.pdf", dpi=300, bbox_inches="tight")
plt.show()

```



Hallmark Enrichment (ORA via Fisher's Exact Test)

```
def run_ora(gene_list, universe, hallmark_sets, min_gs_size=5, p_adj_method="fdr_bh"):
    gene_set = set(gene_list) & universe
    n_uni, n_query = len(universe), len(gene_set)
    rows = []
    for hallmark, gs in hallmark_sets.items():
        gs_in_uni = gs & universe
        if len(gs_in_uni) < min_gs_size:
            continue
        overlap_genes = gene_set & gs_in_uni
        a = len(overlap_genes)
        b = n_query - a
        c = len(gs_in_uni) - a
        d = n_uni - a - b - c
        _, pval = fisher_exact([[a, b], [c, d]], alternative="greater")
        rows.append({"Hallmark": hallmark, "Overlap": a,
                    "GS Size": len(gs_in_uni), "p-value": pval,
                    "Overlap Genes": ", ".join(sorted(overlap_genes))})
    if not rows:
        return pd.DataFrame()
    df = pd.DataFrame(rows)
    _, padj, _, _ = multipletests(df["p-value"], method=p_adj_method)
    df["p.adj"] = padj
    return df.sort_values("p-value").reset_index(drop=True)

hallmark_sets = (
    hallmarks_raw.dropna(subset=["HALLMARK"])
    .query("HALLMARK != ''")
    .groupby("HALLMARK")["GENE_SYMBOL"].apply(set).to_dict()
)
universe = set(ranked_genes["gene_name"])
top100 = set(ranked_genes.head(100)["gene_name"])
consensus = set(ranked_genes.loc[ranked_genes["consensus_significant"], "gene_name"]) \
    if "consensus_significant" in ranked_genes.columns else top100
```

Top 100 Predicted Drivers

```
ora_top100 = run_ora(top100, universe, hallmark_sets)
if ora_top100.empty or (ora_top100["p.adj"] >= 0.05).all():
```

```

    print("No significant hallmark enrichment at p.adj < 0.05 for top 100")
else:
    sig = ora_top100[ora_top100["p.adj"] < 0.05].copy()
    sig["-log10(p.adj)"] = -np.log10(sig["p.adj"])
    fig, ax = plt.subplots(figsize=(7, max(3, len(sig) * 0.45)))
    sns.barplot(data=sig, x="-log10(p.adj)", y="Hallmark", palette="Blues_r", ax=ax)
    ax.axvline(-np.log10(0.05), color="red", linestyle="--", lw=1, label="p.adj = 0.05")
    ax.set_title("Hallmark Enrichment - Top 100 Predicted Drivers")
    ax.legend()
    plt.tight_layout()
    plt.show()
    display(sig[["Hallmark", "Overlap", "GS Size", "p-value", "p.adj"]]
            .style.format({"p-value": "{:.3e}", "p.adj": "{:.3e}"}))

```

C:\Users\siddh\AppData\Local\Temp\ipykernel_43672\2642632014.py:8: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign

```

sns.barplot(data=sig, x="-log10(p.adj)", y="Hallmark", palette="Blues_r", ax=ax)

```




Table 2

	Hallmark	Overlap	GS Size	p-value	p.adj
0	role in cancer	33	343	1.158e-25	1.401e-24
1	function summary	33	344	1.274e-25	1.401e-24
2	escaping programmed cell death	24	218	5.529e-20	4.055e-19
3	types of alteration in cancer	24	231	2.198e-19	1.209e-18
4	genome instability and mutations	18	119	1.211e-17	4.769e-17
5	invasion and metastasis	22	216	1.301e-17	4.769e-17
6	impact of mutation on function	20	189	2.170e-16	6.821e-16
7	differentiation and development	19	167	3.377e-16	9.285e-16
8	proliferative signalling	20	195	4.021e-16	9.830e-16
9	cell division control	17	129	1.148e-15	2.526e-15
10	tumour promoting inflammation	12	61	1.611e-13	3.223e-13
11	cell replicative immortality	10	51	1.949e-11	3.573e-11
12	escaping immune response to cancer	11	70	2.331e-11	3.682e-11
13	angiogenesis	12	91	2.343e-11	3.682e-11
14	senescence	10	63	1.785e-10	2.617e-10
15	fusion partner	12	148	7.134e-09	9.809e-09
16	clinical impact	11	119	7.924e-09	1.025e-08
17	global regulation of gene expression	9	72	1.355e-08	1.657e-08
18	change of cellular energetics	9	85	5.959e-08	6.900e-08
19	suppression of growth	9	132	2.598e-06	2.858e-06
20	mouse model	6	87	1.237e-04	1.296e-04

KEGG Pathway Enrichment

Requires gseapy (pip install gseapy) and internet access for Enrichr API.


```
try:
    import gseapy as gp
    top50_genes = genes.head(50)["gene_name"].tolist()
    enr = gp.enrichr(gene_list=top50_genes, gene_sets=["KEGG_2021_Human"],
                    organism="human", outdir=None, cutoff=0.05)
    results = (enr.results[enr.results["Adjusted P-value"] < 0.05]
              .sort_values("Combined Score", ascending=False).head(15))
    if results.empty:
        results = enr.results.sort_values("Adjusted P-value").head(15)
    fig, ax = plt.subplots(figsize=(8, 6))
```



```

colors = plt.cm.RdYlGn_r(np.linspace(0.1, 0.9, len(results)))
ax.barh(results["Term"].str.replace(r" \(.*\)", "", regex=True),
         results["Combined Score"], color=colors)
sm = plt.cm.ScalarMappable(cmap="RdYlGn_r",
                           norm=plt.Normalize(results["Adjusted P-value"].min(),
                                                results["Adjusted P-value"].max()))
sm.set_array([])
plt.colorbar(sm, ax=ax, shrink=0.6).set_label("Adjusted P-value", fontsize=10)
ax.set_xlabel("Combined Score", fontsize=12)
ax.set_title("KEGG Pathway Enrichment - Top-50 DriverJinn Predictions\n(GGNet,  $\rho=0.2$ )",
             fontsize=12, fontweight="bold")
ax.invert_yaxis()
sns.despine()
plt.tight_layout()
plt.savefig("poster_kegg_enrichment.pdf", dpi=300, bbox_inches="tight")
plt.show()
except ImportError:
    print("gseapy not installed - run: pip install gseapy")

```



Novel Candidate Genes

Top consensus-significant predictions not in any CGC tier:

```
if "consensus_significant" in ranked_genes.columns:
    novel_df = (
        ranked_genes
        .loc[ranked_genes["consensus_significant"] & ~ranked_genes["cgc_any"]]
        [["aggregate_rank", "gene_name", "mean_score", "folds_significant", "total_folds"]]
        .head(20)
    )
    display(novel_df.style.format({"mean_score": "{:.4f}").hide(axis="index"))
else:
    print("consensus_significant column not found - skipping novel candidate table")
```

Table 3

aggregate_rank	gene_name	mean_score	folds_significant	total_folds
42	MNAT1	1.0000	5	5
59	TRAF6	0.9999	5	5
69	MRE11	0.9999	5	5
84	GTF2H1	0.9999	5	5
96	PDS5B	0.9999	5	5
106	DYNC1H1	0.9999	5	5
124	AGO3	0.9999	5	5
133	GTF2H4	0.9999	5	5
147	RPTOR	0.9999	5	5
174	SYNE1	0.9999	5	5
186	ARRB2	0.9999	5	5
188	LRP2	0.9999	5	5
210	YWHAB	0.9999	5	5
211	VCL	0.9999	5	5
215	PAK1	0.9999	5	5
216	MARK3	0.9999	5	5
221	BTRC	0.9999	5	5
239	REV3L	0.9999	5	5
251	SNRPB	0.9999	5	5
257	VWF	0.9999	5	5

Cross-Model Consensus

Combining ranked lists from 6 augmentation strategies to identify genes consistently predicted across strategies.

```
MODEL_LISTS = {
    "random_r0.1"      : "../model_results_2026-03-14/GGNet_random_r0.1_aggregated/GGNet_random_r0.1_aggregated.csv",
    "random_r0.2"      : "../model_results_2026-03-14/GGNet_random_r0.2_aggregated/GGNet_random_r0.2_aggregated.csv",
    "coarsening_r0.1"   : "../model_results_2026-03-15/GGNet_coarsening_r0.1_aggregated/GGNet_coarsening_r0.1_aggregated.csv",
    "coarsening_r0.2"   : "../model_results_2026-03-15/GGNet_coarsening_r0.2_aggregated/GGNet_coarsening_r0.2_aggregated.csv",
    "priority_r0.1"     : "../model_results_2026-03-15/GGNet_priority_r0.1_aggregated/GGNet_priority_r0.1_aggregated.csv",
    "priority_r0.2"     : "../model_results_2026-03-15/GGNet_priority_r0.2_aggregated/GGNet_priority_r0.2_aggregated.csv"
}

dfs = {}
for model, path in MODEL_LISTS.items():
    df = pd.read_csv(path)[["gene_name", "aggregate_rank", "mean_score", "consensus_significant"]]
    df = df.rename(columns={"aggregate_rank": f"rank_{model}",
                           "mean_score": f"score_{model}",
                           "consensus_significant": f"consensus_{model}"})
    dfs[model] = df

merged_models = reduce(lambda l, r: pd.merge(l, r, on="gene_name", how="outer"), dfs.values())

rank_cols = [f"rank_{m}" for m in MODEL_LISTS]
score_cols = [f"score_{m}" for m in MODEL_LISTS]
cons_cols = [f"consensus_{m}" for m in MODEL_LISTS]


merged_models["mean_rank"] = merged_models[rank_cols].mean(axis=1)
merged_models["mean_score_all"] = merged_models[score_cols].mean(axis=1)
merged_models["models_consensus"] = merged_models[cons_cols].sum(axis=1).astype(int)
merged_models["cgc_tier"] = merged_models["gene_name"].map(
    lambda g: "Tier 1" if g in cgc_tier1 else ("Tier 2" if g in cgc_tier2 else "Novel")
)
cross_consensus = merged_models.sort_values("mean_rank").reset_index(drop=True)

print(f"Total genes across all models : {len(cross_consensus)}")
print(f"Consensus in 3/6 models : {(cross_consensus['models_consensus'] >= 3).sum()}")
print(f"Consensus in 4/6 models : {(cross_consensus['models_consensus'] >= 4).sum()}")
print(f"Consensus in all 6 models : {(cross_consensus['models_consensus'] == 6).sum()}")
```

Total genes across all models : 11183
Consensus in 3/6 models : 74
Consensus in 4/6 models : 6
Consensus in all 6 models : 0

Rank Heatmap — Top 50 by Mean Rank

```
top_n    = 50
top_ids = cross_consensus.head(top_n)["gene_name"].tolist()
heat_df = cross_consensus.set_index("gene_name").loc[top_ids, rank_cols].copy()
heat_df.columns = list(MODEL_LISTS.keys())
fig, ax = plt.subplots(figsize=(10, max(6, top_n * 0.28)))
sns.heatmap(heat_df, cmap="RdYlGn_r", ax=ax, linewidths=0.3, linecolor="white",
            cbar_kws={"label": "Aggregate Rank", "shrink": 0.6})
ax.set_xlabel("Augmentation Strategy")
ax.set_title(f"Rank Heatmap - Top {top_n} Genes by Mean Rank Across All Models")
ax.set_xticklabels(ax.get_xticklabels(), rotation=30, ha="right")
plt.tight_layout()
plt.show()
```



Top Cross-Model Consensus Genes (4/6 models)

```
cols_display = ["gene_name", "mean_rank", "mean_score_all", "models_consensus", "cgc_tier"]
fmt = {"mean_rank": "{:.1f}", "mean_score_all": "{:.4f}"}
fmt.update({c: "{:.0f}" for c in rank_cols})
strong_4 = cross_consensus[cross_consensus["models_consensus"] >= 4][cols_display].reset_index()
strong_4.index += 1
strong_4.style.format(fmt).background_gradient(subset=rank_cols, cmap="RdYlGn_r", axis=None)
```

Table

	gene_name	mean_rank	mean_score_all	models_consensus	cgc_tier	rank_random_r0.1	rank_r
1	TRERF1	439.5	0.6662	4	Novel	240	436
2	CBFA2T2	489.3	0.6862	4	Novel	441	693
3	ANAPC2	998.0	0.6053	4	Novel	1322	798
4	NPAS2	1112.3	0.5251	4	Novel	974	1190
5	SLC16A1	1419.7	0.4960	4	Novel	1569	1545
6	TUB	4395.7	0.2400	4	Novel	2837	4021

PPI Network Analysis


Degree Centrality vs Model Score

```
ranked_genes["degree"] = ranked_genes["gene_name"].map(
    lambda g: degree_map.get(name_to_idx.get(g), np.nan)
)

fig, ax = plt.subplots(figsize=(7, 5))
for tier, grp in ranked_genes.dropna(subset=["degree"]).groupby("cgc_label"):
    ax.scatter(grp["degree"], grp["mean_score"], alpha=0.35, s=12,
               color=tier_colors[tier], label=tier, rasterized=True)
ax.set_xlabel("Node Degree (PPI)")
ax.set_ylabel("Mean Model Score")
ax.set_title("Model Score vs Network Degree")
ax.legend(title="CGC Tier")
plt.tight_layout()
```

```
plt.show()

valid = ranked_genes.dropna(subset=["degree"])
rho, pval = spearmanr(valid["degree"], valid["mean_score"])
print(f"Spearman rho (degree vs score): {rho:.3f} p={pval:.2e}")
```



Spearman rho (degree vs score): 0.476 p=0.00e+00

Degree Distribution: Top-100 vs CGC vs All Genes


```
all_degrees      = [degree_map.get(name_to_idx.get(g), 0) for g in ranked_genes["gene_name"] if g != ""]
top100_degrees   = [degree_map.get(name_to_idx.get(g), 0) for g in ranked_genes.head(100)["gene_name"] if g != ""]
cgc1_degrees     = [degree_map.get(name_to_idx.get(g), 0) for g in ranked_genes[ranked_genes["cgc_tier"] == 1]]

fig, ax = plt.subplots(figsize=(7, 4))
```

```

bins = np.linspace(0, max(all_degrees), 60)
ax.hist(all_degrees, bins=bins, alpha=0.4, label=f"All genes (n={len(all_degrees)})",
ax.hist(cgc1_degrees, bins=bins, alpha=0.6, label=f"CGC Tier 1 (n={len(cgc1_degrees)})",
ax.hist(top100_degrees, bins=bins, alpha=0.7, label=f"Top-100 predicted (n={len(top100_degrees)})",
ax.set_xlabel("Node Degree")
ax.set_ylabel("Density")
ax.set_title("Degree Distribution Comparison")
ax.legend()
plt.tight_layout()
plt.show()
print(f"Median degree - All: {np.median(all_degrees):.0f} | "
      f"CGC Tier 1: {np.median(cgc1_degrees):.0f} | "
      f"Top-100: {np.median(top100_degrees):.0f}")

```



Median degree - All: 61 | CGC Tier 1: 126 | Top-100: 177

PPI Subnetwork of Top-50 Predicted Genes

```

top50_nodes = [name_to_idx[g] for g in genes.head(50)["gene_name"] if g in name_to_idx]
G_sub = G_full.subgraph(top50_nodes).copy()

```



```

G_named = nx.relabel_nodes(G_sub, id_to_name)


score_map_50 = genes.set_index("gene_name")["mean_score"].to_dict()
tier_map_50 = genes.set_index("gene_name")["cgc_label"].to_dict()
node_colors50 = [tier_colors.get(tier_map_50.get(n, "Novel"), "#1f77b4") for n in G_named.nodes()]
node_sizes50 = [score_map_50.get(n, 0.5) * 800 for n in G_named.nodes()]

pos = nx.spring_layout(G_named, seed=42, k=1.5)
fig, ax = plt.subplots(figsize=(12, 10))
nx.draw_networkx_edges(G_named, pos, ax=ax, alpha=0.3, edge_color="grey", width=0.8)
nx.draw_networkx_nodes(G_named, pos, ax=ax, node_color=node_colors50,
                        node_size=node_sizes50, edgecolors="white", linewidths=0.8)
nx.draw_networkx_labels(G_named, pos, ax=ax, font_size=7.5, font_weight="bold")
ax.legend(handles=[Patch(color=c, label=t) for t, c in tier_colors.items()],
          title="CGC Tier", loc="upper left")
ax.set_title("PPI Subnetwork - Top 50 Predicted Driver Genes\n"
            "(node size = model score, color = CGC tier)", pad=12)
ax.axis("off")
plt.tight_layout()
plt.show()
print(f"Subgraph density: {nx.density(G_named):.4f}")
connected = list(nx.connected_components(G_named))
print(f"Connected components: {len(connected)} (largest: {max(len(c) for c in connected)} nodes)")

```

CGC Tier

	Tier 1
	Tier 2
	Novel



Cross-Network Consensus Table

```
ggnet      = pd.read_csv("../model_results_2026-03-14/GGNet_random_r0.2_aggregated/"
                        "GGNet_random_r0.2_aggregated_all_genes.csv")
pathnet    = pd.read_csv("../model_results_2026-03-20/PathNet_random_r0.2_aggregated/"
                        "PathNet_random_r0.2_aggregated_all_genes.csv")
```

```

ppnet = pd.read_csv("../model_results_2026-03-20/PPNet_random_r0.2_aggregated/"
                    "PPNet_random_r0.2_aggregated_all_genes.csv")

def cgc_label(gene):
    if gene in cgc_tier1: return "Tier 1"
    if gene in cgc_tier2: return "Tier 2"
    return "Novel"

merged_net = (
    ggnet[["gene_name", "aggregate_rank", "mean_score"]]
    .rename(columns={"aggregate_rank": "GGNet_rank", "mean_score": "GGNet_score"})
    .merge(pathnet[["gene_name", "aggregate_rank", "mean_score"]]
          .rename(columns={"aggregate_rank": "PathNet_rank", "mean_score": "PathNet_score"}),
          on="gene_name", how="inner")
    .merge(ppnet[["gene_name", "aggregate_rank", "mean_score"]]
          .rename(columns={"aggregate_rank": "PPNet_rank", "mean_score": "PPNet_score"}),
          on="gene_name", how="inner")
)
merged_net["CGC"] = merged_net["gene_name"].apply(cgc_label)
merged_net["consensus_rank"] = merged_net[["GGNet_rank", "PathNet_rank", "PPNet_rank"]].mean
merged_net = merged_net.sort_values("consensus_rank")

top20 = merged_net.head(20).copy()
top20.insert(0, "Consensus Rank", range(1, 21))
top20 = top20.rename(columns={"gene_name": "Gene", "GGNet_rank": "GGNet",
                             "PathNet_rank": "PathNet", "PPNet_rank": "PPNet",
                             "CGC": "CGC Status"})
display_cols = ["Consensus Rank", "Gene", "GGNet", "PathNet", "PPNet", "CGC Status"]

ggnet_scores = ggnet.set_index("gene_name")["mean_score"].to_dict()

def style_row(row):
    if row["CGC Status"] == "Novel": return ["background-color: #FFF3CD; font-weight: bold"]
    if row["CGC Status"] == "Tier 1": return ["background-color: #FDECEA"] * len(row)
    return [""] * len(row)

styled = (
    top20[display_cols].style
    .apply(style_row, axis=1)
    .format({"GGNet": "{:.0f}", "PathNet": "{:.0f}", "PPNet": "{:.0f}"})
    .set_caption("Top 20 Consensus Cancer Driver Gene Predictions Across Networks "
                 "(DriverJinn, =0.2) - Novel candidates highlighted in gold")

```

```

.set_table_styles([
    {"selector": "caption", "props": [("font-size", "13px"), ("font-weight", "bold"),
    ("text-align", "left"), ("padding-bottom", "8px")]},
    {"selector": "th", "props": [("background-color", "#2C3E50"), ("color", "white"),
    ("font-size", "11px"), ("text-align", "center"),
    ("padding", "6px 10px")]},
    {"selector": "td", "props": [("font-size", "11px"), ("text-align", "center"),
    ("padding", "5px 10px")]},
    {"selector": "tr:hover td", "props": [("background-color", "#EBF5FB)]},
])
)

novel_top20 = top20[top20["CGC Status"] == "Novel"]["Gene"].tolist()
print(f"Novel candidates in top 20: {novel_top20}")
display(styled)

import dataframe_image as dfi
dfi.export(styled, "novel.png", table_conversion="matplotlib", dpi=200)

```

Novel candidates in top 20: ['SIN3A', 'HDAC2', 'SMARCC1', 'KMT2E', 'SMARCC2', 'SMARCA2', 'RBBP4', 'EED', 'FRS2', 'DOCK7', 'NCOA3']

Table 5: Top 20 Consensus Cancer Driver Gene Predictions Across Networks (DriverJinn, $\alpha=0.2$) — Novel candidates highlighted in gold

	Consensus Rank	Gene	GGNet	PathNet	PPNet	CGC Status
30	1	SIN3A	31	20	118	Novel
22	2	HDAC2	23	82	69	Novel
113	3	NCOR2	121	46	91	Tier 1
55	4	SMARCC1	58	149	85	Novel
224	5	KMT2E	255	22	23	Novel
17	6	SOS1	18	91	193	Tier 1
20	7	CBFB	21	84	209	Tier 1
93	8	KDR	100	113	103	Tier 1
85	9	SMARCC2	92	29	196	Novel
171	10	SMARCA2	190	68	61	Novel
12	11	RBBP4	13	47	294	Novel
196	12	EED	218	34	115	Tier 2
82	13	FRS2	88	132	147	Novel
254	14	DOCK7	290	37	64	Novel
280	15	NCOA3	319	4	80	Novel

	Consensus Rank	Gene	GGNet	PathNet	PPNet	CGC Status
120	16	MED1	129	108	170	Novel
76	17	CHD4	82	19	315	Tier 1
95	18	SHC1	102	281	36	Novel
195	19	MARK3	216	107	100	Novel
2	20	HDAC1	3	122	301	Novel

PPI Neighborhood Subgraph (Consensus Top 20)

Top 20 consensus genes + bridge neighbors connecting 6 of them.

```

target_genes = top20["Gene"].tolist()
target_ids   = [name_to_id[g] for g in target_genes if g in name_to_id]
found_genes  = [g for g in target_genes if g in name_to_id]
target_id_set = set(target_ids)

# 1-hop neighbors connecting 6 target genes
neighbor_target_count = {}
for nid in target_ids:
    for nb in G_full.neighbors(nid):
        if nb not in target_id_set:
            neighbor_target_count[nb] = neighbor_target_count.get(nb, 0) + 1
filtered_neighbors = {nb for nb, cnt in neighbor_target_count.items() if cnt >= 6}
neighbor_ids = target_id_set | filtered_neighbors

H_raw = G_full.subgraph(neighbor_ids).copy()
H      = nx.relabel_nodes(H_raw, id_to_name)

target_set    = set(found_genes)
cgc_color_map = {"Tier 1": "#ca301f", "Tier 2": "#E8A838", "Novel": "#4A90D9"}
cgc_lookup    = top20.set_index("Gene")["CGC Status"].to_dict()

# Edges
target_edges = [(u, v) for u, v in H.edges() if u in target_set and v in target_set]
spoke_edges  = [(u, v) for u, v in H.edges() if (u in target_set) != (v in target_set)]

fig, ax = plt.subplots(figsize=(16, 12))
pos = nx.spring_layout(H, seed=42, k=1.5)

```

```

nx.draw_networkx_edges(H, pos, edgelist=spoke_edges, ax=ax, alpha=0.10, edge_color="#AAAAAA")
nx.draw_networkx_edges(H, pos, edgelist=target_edges, ax=ax, alpha=0.60, edge_color="#555555")

neighbor_nodes = [g for g in H.nodes() if g not in target_set]
nx.draw_networkx_nodes(H, pos, nodelist=neighbor_nodes, ax=ax,
                        node_color="#D5D8DC", node_size=120,
                        edgecolors="#AAAAAA", linewidths=0.8, alpha=0.55)
target_in_H = list(target_set & set(H.nodes()))
nx.draw_networkx_nodes(H, pos, nodelist=target_in_H, ax=ax,
                        node_color=[cgc_color_map.get(cgc_lookup.get(g, "Novel"), "#4A90D9")
                                   for g in target_in_H],
                        node_size=[max(ggnet_scores.get(g, 0.5), 0.1) * 2500 for g in target_in_H],
                        edgecolors="white", linewidths=2.0, alpha=0.95)
nx.draw_networkx_labels(H, pos, labels={g: g for g in target_set if g in H.nodes()},
                        ax=ax, font_size=8, font_weight="bold", font_color="#1A1A1A")

patches = [mpatches.Patch(color=c, label=l) for l, c in cgc_color_map.items()]
patches.append(mpatches.Patch(color="#D5D8DC", label="PPI Neighbor"))
ax.legend(handles=patches, title="CGC Tier", fontsize=11, title_fontsize=11,
          loc="upper left", framealpha=0.9)
ax.set_title(f"PPI Neighborhood - Top 20 Consensus Driver Genes + Bridge Neighbors\n"
             f"(node size = GGNet score, color = CGC tier, "
             f"n={H.number_of_nodes()} nodes, {H.number_of_edges()} edges)",
             fontsize=13, fontweight="bold", pad=14)
ax.axis("off")
plt.tight_layout()
plt.savefig("consensus_subnetwork.pdf", dpi=300, bbox_inches="tight")
plt.savefig("consensus_subnetwork.png", dpi=200, bbox_inches="tight")
plt.show()
print(f"Target genes in graph : {sorted(target_set & set(H.nodes()))}")
print(f"Total nodes          : {H.number_of_nodes()}")
print(f"Total edges           : {H.number_of_edges()}")

```

